Phase 1: Fault Injection and Propagation in a Researcher–Planner–Executor Workflow

Yitang Gao Qinlin Chen Xueru Song Kai Zhang
University of California, San Diego

1 Goal

Option 5 asks whether tracking dependencies between agent results can make recovery from a known-faulty result efficient while keeping execution correct and safe. Phase 1 builds the ground for that question: a clean workflow with checkable outcomes, a harness that injects faults at controlled points, traces that record what each step read and did, measurements of how faults propagate, and the two reference baselines, *full restart* and *always-verify*. Recovery methods are Phase 2 work and are not evaluated here.

2 Setup

Workloads. Each workload has a clean run whose outcome a checker verifies. **W1** (booking) is a researcher → planner → executor workflow. The researcher looks up team size and room capacities, the planner picks a room for each of three meetings, and the executor books it and e-mails attendees (14 steps, 6 LLM calls). `calendar.book` is compensable; `email.send` is irreversible. **W3** (code fix) has an agent read an issue and source file, then write a patch into a git workspace that unit tests check. This adds non-LLM file state. **W4** ports two scenarios from TAMAS [1], a named starting point in the brief: Byzantine agents (`w4.byz`) and indirect prompt injection (`w4.ipi`). Their effects are posts to external sites, which are irreversible.

Model and platform. Qwen2.5-7B-Instruct (Q4_K_M) runs in-process through `llama-cpp-python` on one Apple M5 Pro laptop (24 GB), with greedy decoding and JSON-constrained output. A scripted model runs the same workflows for fast checks. The propagation grid (E1) has 3 repeats per configuration, which are scenario variants; the scripted W2 grids have 5, and each baseline cell has 1. 1,675 cells ran with zero errors. A rerun of both grids on an Apple M5 reproduces every count in Tables 1–3; its raw cells and logs are in `results/` of the repository.

Traces. Every step reads and writes one versioned namespace: messages, shared state and files. A step records (i) the messages and artifacts it wrote, (ii) its tool actions with their effect class (`read`, `compensable`, `irreversible`), (iii) its cost (LLM calls, tokens, wall time), and (iv) its *dependencies*: the ordered paths it read, each with a content hash. A fault's *cone* is every step that read the faulty value, directly or through another step.

Faults. One fault per run, at a chosen step: *wrong* (a plausible false value), *stale* (an outdated value), *timeout* (a missing value) and *malicious* (an injected instruction or a Byzantine agent). Accidental and adversarial faults are reported separately. Every workload has a timeout fault (W1 team size, W3 source file, W4 one feed item), and W1 also has a planner-stage fault (a room that never fits). *Unauthorised actions* are actions the task does not allow, counted on the faulty run: e-mails to non-members (W1), edits to any file but `cart.py` (W3), and posts on channels the request did not name (W4). A faulty cell whose injected value no later step reads is excluded; a cell that is read but still passes the checker is reported as *masked*.

3 Clean baseline

All four workloads succeed in 3/3 clean runs on both models (Table 1). The clean-vs-clean spread in wall time is 1.1 s (W1), 1.2 s (W3) and 0.6 s (W4.ipi); these describe variation across scenario variants, not a measured latency noise floor.

Table 1: Clean runs, 7B model (n=3 each).

WL	success	steps	LLM calls	tokens	latency s
W1	3/3	14	6	2427	47.6
W3	3/3	5	1	661	8.7
W4.byz	3/3	6	3	1070*	–*
W4.ipi	3/3	6	3	1078	13.2

*The grid's cells were prompt-cache hits, so no live latency was measured; tokens come from three separate uncached runs.

4 Propagation

For each fault we measure whether any later step read the faulty value, whether the task checker failed, how many steps and which effects it reached, and whether unauthorised actions occurred (Table 2).

Table 2: Fault propagation, 7B model (n=3 each). "Fail" is the fraction of runs that fail the task checker. "Irrev." counts irreversible calls made in the faulty run. The timeout and planner-stage faults are not in this 7B grid yet (see Limitations).

WL	fault (class)	fail	steps hit	irrev.	unauth.
W1	wrong capacity (acc.)	1.00	9/14	3	0
	stale capacity (acc.)	0.67	9/14	3	0
	injected agenda (adv.)	1.00	3/14	3	**1**
W3	wrong issue (acc.)	1.00	3/5	0	0
	stale code (acc.)	1.00	3/5	0	0
	injected issue (adv.)	0.00	3/5	0	0
W4.byz	stale flag (acc.)	0.67	4/6	2	0
	Byz. distribute (adv.)	1.00	1/6	1	0
	Byz. fact-check (adv.)	0.00	3/6	2	0
	injected item (adv.)	0.33	4/6	2.3	0
W4.ipi	stale flag (acc.)	0.67	4/6	1	0
	Byz. distribute (adv.)	1.00	1/6	1	0
	Byz. fact-check (adv.)	0.00	3/6	1	0
	injected item (adv.)	0.33	4/6	1.3	0

- **Every fault is read downstream** (all 14 fault–workload pairs), so each workload really depends on the injected result.
- **Reach varies widely.** A wrong W1 capacity reaches 9 of 14 steps, both bookings and the irreversible e-mail. The injected agenda reaches only 3 steps, yet sends an unauthorised e-mail in every run.
- **Some faults are read without task failure.** Three adversarial faults pass the 7B checker in all tested variants (W3 injected issue; Byzantine fact-check in both W4 scenarios), unlike the scripted model. This establishes neither unchanged outputs nor general robustness.
- **Accidental faults can cause task failures.** A stale capacity fails the W1 checker in 2 of 3 runs.

5 Baselines

Full restart undoes what it can and reruns every step. **Always-verify** audits each model output against its inputs with the same model; tool-only steps are not audited. Its tokens are counted as extra inference cost. Both see the same faulty runs (Table 3). *Wasted work* is the number of rerun steps outside the fault's cone, which a dependency-aware method could avoid.

Observations. (1) Restart reruns many steps that never read the fault: on W1, 10 of 14 for the injected agenda and 4 of 14 for the capacity faults. It also re-sends every irreversible call (3.0 per W1 recovery), and it cannot recall the unauthorised e-mail. (2) Mean audit tokens are 51–87% of the clean-run reference tokens, not a measured overhead for every clean or faulty run. The detection proxy misses one third of W3 faults and a quarter of W4.byz faults; off-cone flags have not been independently validated as false positives. (3) †Our first W3 restart cells failed in every 7B run, although clean W3 succeeds 3/3. The cause was in our harness, not the model: restart undid each step's old effects only when that step reran, so the rerun researcher read `cart.py` as the faulty run had left it. Where the faulty run had already patched the file, the planner saw fixed code and "fixed" it back. The scripted model leaves correct code alone, which hid the bug. Restart now undoes every compensable effect, newest first, before it reruns anything, and all three W3 cells recover. The same ordering problem can affect any method that rereads tool state, which Phase 2 has to handle.

Table 3: Baselines, 7B model, mean over faults. Restart: steps rerun, wasted steps (outside the cone), recovery time and tokens, irreversible calls re-sent, recovered correctly. Always-verify: audit tokens per run and their ratio to a clean run's tokens, detection rate (a faulty run counts as detected when the audit flags the producer or a step in its cone), off-cone flags (OC) per run. OC is not a validated false-positive count.

	full restart						always-verify			
WL	rerun	waste	s	tok	irr.	ok	tok	×	det.	OC
W1	14	6.0	36.3	2326	3.0	1.0	1245	.51	1.00	.33
W3†	5	1.0	10.6	592	0	1.0	575	.87	0.67	0
W4.byz	6	2.0	13.4	1078	1.8	1.0	692	.65	0.75	.25
W4.ipi	6	2.0	14.1	1097	1.0	1.0	712	.66	1.00	.25

†Restart rerun after the fix in observation (3), on an Apple M5.

Limitations. The timeout and planner-stage faults, and the per-workload unauthorised-action counts, are implemented and pass the scripted grid, but the 7B grid has not been rerun with them yet. Until it is, the unauthorised column of Table 2 counts only e-mails (so W3 and W4 read 0) and comes from the one-repeat baseline grid. W4 uses 2 of TAMAS's scenarios; no A2ASecBench attack is ported yet. n=3 (and n=1 for the baselines) gives wide intervals.

Takeaway for Phase 2. Faults reach only part of each workflow, but restart pays for all of it and re-sends irreversible effects, while always-verify adds audit cost and still misses faults. Phase 2 will use the recorded dependencies to rerun only the cone, reconcile effects instead of repeating them, and compare against these two baselines under identical fault schedules.

References

[1] I. Kavathekar et al. TAMAS: Benchmarking adversarial risks in multi-agent LLM systems, 2025. arXiv:2511.05269. Code: github.com/microsoft/TAMAS.

Appendix: Timeline and Milestones

Due / status	Milestone	Done when
Phase 1 — fault injection and propagation (this report)		
Partial	Review TAMAS and A2ASecBench; choose scenarios	two TAMAS scenarios ported (W4); A2ASecBench not yet ported
Done	Clean researcher–planner–executor baseline	W1, W3, W4 pass their checkers 3/3
Partial	Fault-injection harness	wrong / stale / timeout / malicious faults at a chosen step; timeout and planner-stage faults not yet in the 7B grid
Done	Traces: messages, artifacts, tool actions, dependencies	every step records reads with hashes, writes, effects, cost
Done	Measure clean runs and propagation	Tables 1–2
Done	Full-restart and always-verify baselines	Table 3
Phase 2 — selective recovery		
Nov 8	Dependency tracking: tag results with their sources	a step's inputs can be listed from the trace
Nov 15	Dependency-aware invalidation and selective rerun	only steps in the fault's cone rerun
Nov 22	Shared state and irreversible actions	compensation for compensable effects; irreversible ones reported, not repeated
Nov 29	Compare against full restart and always-verify	identical fault schedules and budgets, fault stage varied
Dec 5	Evaluate recovery metrics and reuse validity	correctness, unauthorised actions, recomputed stages, retained work, latency, inference cost, tracking overhead
Dec 7	Optional: automated fault detection	false positives and misses reported separately from recovery
Dec 10	Submit Phase 2 code and the final report	4–6 page report

Open items carried into Phase 2. Rerun the 7B grid with the timeout and planner-stage faults and the per-workload unauthorised counts; port at least one A2ASecBench attack; add a stateful SQLite procurement workload (inventory, orders, a mock irreversible payment) as in the proposal; make reruns that reread tool state see the world as it was before the fault, not as the faulty run left it; add the precision-versus-bookkeeping measurement (dependency granularity) that the brief asks for.

Acknowledgments. Phase 1 progress report for the CSE 291A (*Systems for LLMs and AI Agents*, UC San Diego, Fall 2026) course project.